‹ Return to Review

Attempt 1

All Questions

Question 1: Skipped

There is a requirement for a data store in AWS. Below are the requirements for the data storea) Ability to perform SQL queriesb) Integration with existing business intelligence toolsc) High concurrency workload that generally involves reading and writing all of the columns for a small number of records at a timeWhich of the following would be the ideal data store that can be used for such requirements. Choose 2 answers from the options below

- A. AWS Redshift
- B. AWS RDS **(Correct)**
- C. AWS Aurora **(Correct)**
- D. AWS S3

Explanation

The AWS Documentation mentions this as a best practice. Because Amazon Redshift is a SQL-based relational database management system (RDBMS), it is compatible with other RDBMS applications and business intelligence tools. Although Amazon Redshift provides the functionality of a typical RDBMS, including online transaction processing (OLTP) functions, it is not designed for these workloads. If you expect a high concurrency workload that generally involves reading and writing all of the columns for a small number of records at a time you should instead consider using Amazon RDS or Amazon DynamoDB For more information on AWS Cloud best practises, please visit the following URL:
https://d0.awsstatic.com/whitepapers/AWS_Cloud_Best_Practices.pdf

Question 2: Skipped

A company is currently using Redshift in AWS. There is a mandate that the Redshift cluster is used in a cost effective manner. As an architect which of the following should be consider to ensure cost effectiveness.

- A. Use Spot instances for the underlying nodes in the cluster
- B. Ensure that unnecessary manual snapshots of the cluster are deleted. **(Correct)**
- C. Ensure VPC Enhanced Routing is enabled

D. Ensure that Cloudwatch metrics are disabled

Explanation

The AWS Documentation mentions the following Amazon Redshift provides free storage for snapshots that is equal to the storage capacity of your cluster until you delete the cluster. After you reach the free snapshot storage limit, you are charged for any additional storage at the normal rate. Because of this, you should evaluate how many days you need to keep automated snapshots and configure their retention period accordingly, and delete any manual snapshots that you no longer need. For more information on working with Redshift snapshots, please visit the following URL: https://docs.aws.amazon.com/redshift/latest/mgmt/working-with-snapshots.html

Question 3: Skipped

A company has a set of resources hosted in a VPC. They have acquired another company and they have their own set of resources hosted in AWS. The requirement now is to ensure that resources in the VPC of the parent company can access the resources in the VPC of the child company. How can this be accomplished.

A. Establish a NAT instance to establish communication across VPC's

B. Establish a NAT gateway to establish communication across VPC's

C. Use a VPN connection to peer both VPC's

D. Use VPC Peering to peer both VPC's (Correct)

Explanation

The AWS Documentation mentions the following about VPC Peering A VPC peering connection is a networking connection between two VPCs that enables you to route traffic between them privately. Instances in either VPC can communicate with each other as if they are within the same network. You can create a VPC peering connection between your own VPCs, with a VPC in another AWS account, or with a VPC in a different AWS Region. For more information on VPC Peering, please visit the following URL: https://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/vpc-peering.html

Question 4: Skipped

An application consists of the following architecture.a. EC2 Instances in a single AZ behind an ELB.b. A NAT instance which is used to ensure that instances can download updates from the internet.Which of the following can be used to ensure better fault tolerance in this setup. Choose 2 answers from the options given below

A. Add more instances in the existing Availability Zone

B. Add an Autoscaling Group to the setup (Correct)

C. Add more instances in another Availability Zone (Correct)

D. Add another ELB for more fault tolerance

Explanation

The AWS Documentation mentions the following Adding Auto Scaling to your application architecture is one way to maximize the benefits of the AWS cloud. When you use Auto Scaling, your applications gain the following benefits: Better fault tolerance. Auto Scaling can detect when an instance is unhealthy, terminate it, and launch an instance to replace it. You can also configure Auto Scaling to use multiple Availability Zones. If one Availability Zone becomes unavailable, Auto Scaling can launch instances in another one to compensate. Better availability. Auto Scaling can help you ensure that your application always has the right amount of capacity to handle the current traffic demands. For more information on the benefits of AutoScaling, please visit the following URL: https://docs.aws.amazon.com/autoscaling/ec2/userguide/auto-scaling-benefits.html

Question 5: Skipped

A company currently has a lot of data hosted on their On-premise infrastructure. They are now running out of storage space and looking towards a quick win solution using AWS. Which of the following would allow them to easily extend their data infrastructure to AWS>

A. Let the company start using Gateway Cached volumes **(Correct)**

B. Let the company start using Gateway Stored volumes

C. Let the company start using the Simple Storage service

D. Let the company start using Amazon Glacier

Explanation

One can easily start using Volume gateways to start storing their data in S3. One can use the Cached volumes. The AWS Documentation mentions the following You store your data in Amazon Simple Storage Service (Amazon S3) and retain a copy of frequently accessed data subsets locally. Cached volumes offer a substantial cost savings on primary storage and minimize the need to scale your storage on-premises. You also retain low-latency access to your frequently accessed data. For more information on Storage gateways please visit the following URL:
https://docs.aws.amazon.com/storagegateway/latest/userguide/WhatIsStorageGateway.html

Question 6: Skipped

Company salespeople upload their sales figures daily. A Solutions Architect needs a durable storage solution for these documents that also protects against users accidentally deleting important documents. Which action will protect against unintended user actions?

A. Store data in an EBS volume and create snapshots once a week.

B. Store data in an S3 bucket and enable versioning. **(Correct)**

C. Store data in two S3 buckets in different AWS regions.

D. Store data on EC2 instance storage.

Explanation

Amazon S3 has the option for versioning as shown below. The versioning is on the bucket level and can be used to recover prior versions of an object. For more information on Amazon S3, please visit the following URL: https://aws.amazon.com/s3/

Question 7: Skipped

An application requires a highly available relational database with an initial storage capacity of 8 TB. The database will grow by 8 GB every day. To support expected traffic, at least eight read replicas will be required to handle database reads. Which option will meet these requirements?

A. DynamoDB

B. Amazon S3

C. Amazon Aurora

D. Amazon Redshift **(Correct)**

Explanation

The AWS Documentation mentions the following Amazon Redshift is a fully managed, petabyte-scale data warehouse service in the cloud. You can start with just a few hundred gigabytes of data and scale to a petabyte or more. This enables you to use your data to acquire new insights for your business and customers. For more information on AWS Redshift, please visit the following URL: https://docs.aws.amazon.com/redshift/latest/mgmt/welcome.html

Question 8: Skipped

A company has an application in which objects from S3 are given to users. Some users across the globe are complaining of slow response times. Which of the following additional steps would allow for a COST effective solution and also ensure that the users get the desired optimal response to objects from S3.

A. Use S3 replication to replicate the objects to regions closest to the users.

B. Ensure S3 transfer acceleration is enabled to ensure all users get the desired response times.

C. Place an ELB in from S3 to distribute the load across S3

D. Place the S3 bucket behind a cloudfront distribution **(Correct)**

Explanation

The AWS Documentation mentions the following If your workload is mainly sending GET requests, in addition to the preceding guidelines, you should consider using Amazon CloudFront for performance optimization. Integrating Amazon CloudFront with Amazon S3, you can distribute content to your users with low latency and a high data transfer rate. You will also send fewer direct requests to Amazon S3, which will reduce your costs. For example, suppose that you have a few objects that are very popular. Amazon CloudFront fetches those objects from Amazon S3 and caches them. Amazon CloudFront can then serve future requests for the objects from its cache, reducing the number of GET requests it sends to Amazon S3. For more information on performance considerations in S3, please visit the following URL: https://docs.aws.amazon.com/AmazonS3/latest/dev/request-rate-perf-considerations.html

Question 9: Skipped

An application needs to have a messaging system in AWS. It is of the uttermost importance that the order of messages is preserved and duplicate messages are not sent. Which of the following services can help fulfil this requirement

- A. AWS SQS **(Correct)**
- B. AWS SNS
- C. AWS Config
- D. AWS ELB

Explanation

One can use SQS FIFO queues for this purpose. The AWS Documentation mentions the following on SQS FIFO Queues. Amazon SQS is a reliable and highly-scalable managed message queue service for storing messages in transit between application components. FIFO queues complement the existing Amazon SQS standard queues, which offer high throughput, best-effort ordering, and at-least-once delivery. FIFO queues have essentially the same features as standard queues, but provide the added benefits of supporting ordering and exactly-once processing. FIFO queues provide additional features that help prevent unintentional duplicates from being sent by message producers or from being received by message consumers. Additionally, message groups allow multiple separate ordered message streams within the same queue. For more information on SQS FIFO Queues, please visit the following URL: https://aws.amazon.com/about-aws/whats-new/2016/11/amazon-sqs-introduces-fifo-queues-with-exactly-once-processing-and-lower-prices-for-standard-queues/

Question 10: Skipped

A company is planning on building an application using the services available on AWS. The application will be stateless in nature. Which of the following would be an ideal compute service which can be used. The service should have the ability to scale accordingly

- A. AWS DynamoDB
- B. AWS Lambda **(Correct)**
- C. AWS S3

D. AWS SQS

Explanation

The following is mentioned in the AWS Whitepaper which supplements the ability to use AWS Lambda for this requirement A stateless application is an application that needs no knowledge of previous interactions and stores no session information. Such an example could be an application that, given the same input, provides the same response to any end user. A stateless application can scale horizontally since any request can be serviced by any of the available compute resources (e.g., EC2 instances, AWS Lambda functions) For more information on AWS Cloud best practices, please visit the following URL: https://d1.awsstatic.com/whitepapers/AWS_Cloud_Best_Practices.pdf

Question 11: Skipped

A company has a set of EC2 Instances hosted on the AWS Cloud. These instances form a web server farm which services a web application that is accessed by users on the internet. Which of the following would help make this architecture more fault tolerant. Choose 2 answers from the options given below

A. Ensure the Instances are placed in separate Availability Zones **(Correct)**

B. Ensure the Instances are placed in separate Regions

C. Use an AWS Load Balancer to distribute the traffic **(Correct)**

D. Use Autoscaling to distribute the traffic

Explanation

The AWS Documentation mentions the following A load balancer distributes incoming application traffic across multiple EC2 instances in multiple Availability Zones. This increases the fault tolerance of your applications. Elastic Load Balancing detects unhealthy instances and routes traffic only to healthy instances. For more information on the AWS Classic Load balancer, please visit the following URL: https://docs.aws.amazon.com/elasticloadbalancing/latest/classic/introduction.html

Question 12: Skipped

You are planning on hosting an application on Ec2 Instances that will be used to process logs. This application is not that critical and can resume even after an interruption. Which of the following steps can help provide a COST effective solution.

A. Ensure to use Reserved Instances for the underlying EC2 Instances

B. Ensure to use Provisioned IOPS for the underlying EBS volumes

C. Ensure to use Spot Instances for the underlying EC2 Instances **(Correct)**

D. Ensure to use S3 as the underlying data layer

Explanation

One effective solution would be to use Spot Instances in this scenario. Additionally the AWS Documentation mentions the following Spot Instances are a cost-effective choice if you can be flexible about when your applications run and if your applications can be interrupted. For example, Spot Instances are well-suited for data analysis, batch jobs, background processing, and optional tasks. For more information on using Spot Instances, please visit the following URL: https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/using-spot-instances.html

Question 13: Skipped

A company stores their log data in an S3 bucket. They now need to have search capabilities available for the data in S3. How can this be achieved in an efficient and in on-going manner. Choose 2 answers from the options below. Each answer is part of the solution

- A. Use an AWS Lambda function which gets triggered whenever data is added to the S3 bucket. **(Correct)**
- B. Create a Lifecycle policy for the S3 bucket
- C. Load the data into Amazon ElasticSearch **(Correct)**
- D. Load the data into Glacier

Explanation

AWS Elastic search provides full search capabilities and can be used for the log files stored in the S3 bucket The AWS Documentation mentions the following with regards to the integration of Elastic search with S3 You can integrate your Amazon ES domain with Amazon S3 and AWS Lambda. Any new data sent to an S3 bucket triggers an event notification to Lambda, which then runs your custom Java or Node.js application code. After your application processes the data, it streams the data to your domain. For more information on integration between Elastic Search and S3, please visit the following URL: https://docs.aws.amazon.com/elasticsearch-service/latest/developerguide/es-aws-integrations.html

Question 14: Skipped

A company is planning on deploying a batch processing application in AWS. Which of the following is an ideal way to host this application. Choose 2 answers from the options below. Each answer is part of the solution

- A. Copy the batch processing application to an ECS container
- B. Create a docker image of your batch processing application. **(Correct)**
- C. Deploy the image as an Amazon ECS task **(Correct)**
- D. Deploy the container behind the ELB

Explanation

The AWS Documentation mentions the following Docker containers are particularly suited for batch job workloads. Batch jobs are often short-lived and embarrassingly parallel. You can package your batch processing application into a Docker image so that you can deploy it anywhere, such as in an Amazon ECS task For more information on the use cases for AWS ECS, please visit the following URL: https://docs.aws.amazon.com/AmazonECS/latest/developerguide/common_use_cases.html

Question 15: Skipped

An architecture consists of the followinga) A primary and secondary infrastructure hosted in AWS.b) Both infrastructures consists of ELB, Autoscaling and EC2 resourcesHow should Route53 be configured to ensure proper failover incase the primary infrastructure goes down.

- A. Configure a primary routing policy
- B. Configure a weighted routing policy
- C. Configure a Multi-Answer routing policy
- D. Configure a failover routing policy **(Correct)**

Explanation

The AWS Documentation mentions the following You can create an active-passive failover configuration by using failover records. You create a primary and a secondary failover record that have the same name and type, and you associate a health check with each. For more information on DNS failover using Route53, please visit the following URL: https://docs.aws.amazon.com/Route53/latest/DeveloperGuide/dns-failover-configuring-options.html

Question 16: Skipped

A company has a requirement to provision test environments quickly. They want to have the ability to tear them down also easily to ensure cost is optimized. How can this be achieved.

- A. Use Cloudformation templates to provision the resources accordingly. **(Correct)**
- B. Use a custom script to create and tear down the resources
- C. Use IAM policies to have a policy to provision the resources and tear them down accordingly.
- D. Use Autoscaling groups to provision the resources on demand.

Explanation

The Cost optimization whitepaper from AWS mentions the following AWS CloudFormation provides templates that you can use to create AWS resources and provision them in an orderly and predictable fashion. This can be useful for creating short-lived environments, such as test environments For more information on the whitepaper, please visit the following URL: https://d1.awsstatic.com/whitepapers/architecture/AWS-Cost-Optimization-Pillar.pdf

Question 17: Skipped

A company wants to self-manage a database environment. Which of the following should be adopted to fulfil this requirement

- A. Use the DynamoDB service
- B. Provision the database using the AWS RDS service
- C. Provision the database using the AWS Aurora service
- D. Create an EC2 Instance and install the database service accordingly. **(Correct)**

Explanation

If you want to self-manage a database , then you should have an EC2 Instance and then you will have complete control over the underlying database instance. For more information on Amazon EC2, please visit the following URL: https://aws.amazon.com/ec2/

Question 18: Skipped

A company is migrating an on-premise 5TB MySQL database to AWS. The company expects the database to continue increasing in size. Which Amazon RDS engine meets these requirements?

- A. MySQL
- B. Microsoft SQL Server
- C. Oracle
- D. Amazon Aurora **(Correct)**

Explanation

The AWS Documentation supports the mentioned requirements which is supported by AWS Aurora Amazon Aurora (Aurora) is a fully managed, MySQL- and PostgreSQL-compatible, relational database engine. It combines the speed and reliability of high-end commercial databases with the simplicity and cost-effectiveness of open-source databases. It delivers up to five times the throughput of MySQL and up to three times the throughput of PostgreSQL without requiring changes to most of your existing applications. All Aurora Replicas return the same data for query results with minimal replica lag—usually much less than 100

milliseconds after the primary instance has written an update For more information on AWS Aurora, please visit the following URL: http://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/Aurora.Overview.html

Question 19: Skipped

A company wants to have a 50 Mbps dedicated connection to its AWS resources. Which of the below services can help fulfil this requirement

- A. Virtual private gateway
- B. Virtual private connection
- C. Direct Connect **(Correct)**
- D. Internet gateway

Explanation

AWS Direct Connect makes it easy to establish a dedicated network connection from your premises to AWS. Using AWS Direct Connect, you can establish private connectivity between AWS and your datacenter, office, or colocation environment, which in many cases can reduce your network costs, increase bandwidth throughput, and provide a more consistent network experience than Internet-based connections For more information on AWS direct connect please visit the below URL: https://aws.amazon.com/directconnect/

Question 20: Skipped

You work for a company that stores records for a minimum of 10 years. Most of these records will never be accessed but must be made available upon request (within a few hours). What is the most cost-effective storage option? Choose the correct answer from the options below

- A. Simple Storage Service
- B. EBS Volumes
- C. Glacier **(Correct)**
- D. AWS Import/Export

Explanation

Amazon Glacier is a secure, durable, and extremely low-cost cloud storage service for data archiving and long-term backup. Customers can reliably store large or small amounts of data for as little as $0.004 per gigabyte per month, a significant savings compared to on-premises solutions. To keep costs low yet suitable for varying retrieval needs, Amazon Glacier provides three options for access to archives, from a few minutes to several hours. For more information on Amazon Glacier, please refer to the below link: https://aws.amazon.com/glacier/

Question 21: Skipped

A company wants to store their documents in AWS. Initially these documents will be used frequently. After a duration of 6 months, these documents need to be archived. How would you architect this requirement?

- A. Store the files in Amazon EBS and create a lifecycle policy to remove the files after 6 months.
- B. Store the files in Amazon S3 and create a lifecycle policy to remove the files after 6 months. **(Correct)**
- C. Store the files in Amazon Glacier and create a lifecycle policy to remove the files after 6 months.
- D. Store the files in Amazon EFS and create a lifecycle policy to remove the files after 6 months.

Explanation

The AWS Documentation mentions the following on lifecycle policies Lifecycle configuration enables you to specify the lifecycle management of objects in a bucket. The configuration is a set of one or more rules, where each rule defines an action for Amazon S3 to apply to a group of objects. These actions can be classified as follows: · Transition actions – In which you define when objects transition to another storage class. For example, you may choose to transition objects to the STANDARD_IA (IA, for infrequent access) storage class 30 days after creation, or archive objects to the GLACIER storage class one year after creation. · Expiration actions – In which you specify when the objects expire. Then Amazon S3 deletes the expired objects on your behalf. For more information on AWS S3 Lifecycle policies, please visit the following URL: https://docs.aws.amazon.com/AmazonS3/latest/dev/object-lifecycle-mgmt.html

Question 22: Skipped

When managing permissions for the API gateway, what can be used to ensure that the right level of permissions are given to developers, IT admins and users? These permissions should be easily managed.

- A. Use the secure token service to manage the permissions for the different users
- B. Use IAM Policies to create different policies for the different types of users. **(Correct)**
- C. Use the AWS Config tool to manage the permissions for the different users
- D. Use IAM Access Keys to create sets of keys for the different types of users.

Explanation

The AWS Documentation mentions the following You control access to Amazon API Gateway with IAM permissions by controlling access to the following two API Gateway component processes: · To create, deploy, and manage an API in API Gateway, you must grant the API developer permissions to perform the required actions supported by the API management component of API Gateway. · To call a deployed API or to refresh the API caching, you must grant the API caller permissions to perform required IAM actions supported by the API execution component of API Gateway. For more information on permissions with the

API gateway, please visit the following URL: https://docs.aws.amazon.com/apigateway/latest/developerguide/permissions.html

Question 23: Skipped

Your development team wants to start making use of EC2 Instances to host their application and web servers. In the space of automation, they want the Instances to always download the latest version of the Web and application servers when the Instances are launched. As an architect what would you recommend?

- A. Ask the development team to create scripts which can be added to the User Data section when the instance is launched **(Correct)**
- B. Ask the development team to create scripts which can be added to the Meta Data section when the instance is launched
- C. Use Autoscaling Groups to install the Web and application servers when the instances are launched
- D. Use EC2 Config to install the Web and application servers when the instances are launched

Explanation

The AWS Documentation mentions the following When you launch an instance in Amazon EC2, you have the option of passing user data to the instance that can be used to perform common automated configuration tasks and even run scripts after the instance starts. You can pass two types of user data to Amazon EC2: shell scripts and cloud-init directives. You can also pass this data into the launch wizard as plain text, as a file (this is useful for launching instances using the command line tools), or as base64-encoded text (for API calls). For more information on User Data, please visit the following URL: https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/user-data.html

Question 24: Skipped

Your company has an application that looks into uploading , processing and publishing videos posted by users. They currently have the following architecture for this applicationa) A set of EC2 Instances which look into taking the videos uploaded by users and putting them n S3 bucketsb) A set of EC2 worker processes to process and publish the videosc) An Autoscaling Group for the EC2 worker processesWhich of the following ca be added to the architecture to make it more reliable

- A. Amazon SQS **(Correct)**
- B. Amazon SNS
- C. Amazon Cloudfront
- D. Amazon SES

Explanation

Amazon SQS is used to decouple systems. It can store the requests to process videos which can be picked up by the Worker processes. The AWS Documentation mentions the following Amazon Simple Queue Service (Amazon SQS) offers a reliable, highly-scalable hosted queue for storing messages as they travel between applications or microservices. It moves data between distributed application components and helps you decouple these components For more information on AWS SQS, please visit the following URL: https://docs.aws.amazon.com/AWSSimpleQueueService/latest/SQSDeveloperGuide/Welcome.html

Question 25: Skipped

There is an urgent requirement to monitor some database metrics for a database hosted on AWS and send notifications. Which AWS services can accomplish this? Choose 2 answers from the options given below.

- [ ] A. Amazon Simple Email Service
- [ ] B. Amazon CloudWatch **(Correct)**
- [ ] C. Amazon Simple Queue Service
- [ ] D. Amazon Route 53
- [ ] E. Amazon Simple Notification Service **(Correct)**

Explanation

Amazon Cloudwatch will be used to monitor the IOP's metrics from the RDS instance and Amazon Simple Notification Service will be used to send the notification if any alarm is triggered. For more information on Cloudwatch and SNS, please visit the below URLS: https://aws.amazon.com/cloudwatch/ https://aws.amazon.com/sns/

Question 26: Skipped

You have a business-critical two tier web app currently deployed in 2 availability zones in a single region, using Elastic Load Balancing and Auto-Scaling. The app depends on synchronous replication at the database layer. The application needs to remain fully available even if one application AZ goes off-line and AutoScaling cannot launch new instances in the remaining AZ. How can the current architecture be enhanced to ensure this?

- ( ) A. Deploy in 2 regions using Weighted Round Robin with AutoScaling minimums set of 50% peak load per Region.
- ( ) B. Deploy in 3 AZ with Autoscaling minimum set to handle 33 percent peak load per zone.
- ( ) C. Deploy in 3 AZ with Autoscaling minimum set to handle 50 percent peak load per zone. **(Correct)**
- ( ) D. Deploy in 2 regions using Weighted Round Robin with AutoScaling minimums set of 100% peak load per Region.

Explanation

Since the requirement is that the application should never go down even if an AZ is not available, we need to maintain 100% availability. Option A and D are incorrect because region deployment is not possible for ELB. ELB's can manage traffic within a region and not between regions. Option B is incorrect because even if one AZ goes down, we would be operating at only 66% and not the required 100%. For more information on Autoscaling please visit the below URL:
https://aws.amazon.com/autoscaling/

Question 27: Skipped

You have been tasked with creating a VPC network topology for your company. The VPC network must support both internet-facing applications and internally-facing applications accessed only over VPN. Both Internet-facing and internally-facing applications must be able to leverage at least 3 AZs for high availability. At a minimum, how many subnets must you create within your VPC to accommodate these requirements?

- A. 2
- B. 3
- C. 4
- D. 6 **(Correct)**

Explanation

Since each subnet corresponds to one availability zone and you need 3 AZ's for the internet and intranet applications, hence you need 6 subnets. For more information on VPC and subnets please visit the below URL:
http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_Subnets.html

Question 28: Skipped

You have the following architecture deployed in AWSa) A set of EC2 Instances which sit behind an ELBb) A database hosted in AWS RDSOff late the performance on the database has been slacking due to the high number of read requests. Which of the following can be added to the architecture to alleviate the performance issue.

- A. Enable Multi-AZ to add a secondary read-only DB in another AZ
- B. Use Elastic Cache in front of the database **(Correct)**
- C. Use AWS Cloudfront in front of the database
- D. Use DynamoDB to offload all the reads. Populate the common read items in a separate table.

Explanation

Amazon Elastic Cache is an in-memory cache which can be used to cache common read requests. The diagram from the AWS

Documentations shows how caching can be added to an existing architecture For more information on database caching please visit the below URL: https://aws.amazon.com/caching/database-caching/

Question 29: Skipped

An application is currently hosted on an EC2 Instance which has attached EBS volumes. The data on these volumes is frequently accessed. But after a duration of a week, the documents need to be moved to infrequent access storage. Which of the following would be the ideal EBS volume type to use.

- A. EBS Provisioned IOPS SSD
- B. EBS Throughput Optimized HDD
- C. EBS General Purpose SSD
- D. EBS Cold HDD **(Correct)**

Explanation

The AWS Documentation mentions the following Cold HDD (sc1) volumes provide low-cost magnetic storage that defines performance in terms of throughput rather than IOPS. With a lower throughput limit than st1, sc1 is a good fit ideal for large, sequential cold-data workloads. If you require infrequent access to your data and are looking to save costs, sc1 provides inexpensive block storage. For more information on the various EBS Volume types please visit the below URL: https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/EBSVolumeTypes.html

Question 30: Skipped

A customer wants to import their existing virtual machines to the cloud. Which service can they use for this? Choose one answer from the options given below.

- A. VM Import/Export **(Correct)**
- B. AWS Import/Export
- C. AWS Storage Gateway